

Modelling Altcoin Price Variation with Sentiment Based Predictors

A Thesis
Presented to
Department of Statistical Science
Duke University

Daniel Spottiswood

May 2020

Approved for the
Bachelor of Science in Statistical Science

Sayan Mukherjee

Amy Herring

Fan Li

Dus X. Name, DUS

Acknowledgements

I want to thank a few people.

Preface

This is an example of a thesis setup to use the reed thesis document class (for LaTeX) and the R bookdown package, in general.

Table of Contents

Introduction	3
Literature Review	5
Chapter 1: Data	7
1.1 Twitter	7
1.1.1 Tweet Removal	8
1.1.2 Sentiment Analysis	9
1.2 Google Search Volume	10
1.3 Price and Volume	11
1.4 Exploratory Data Analysis	12
Chapter 2: Methodology	15
2.1 Log Returns	15
2.2 XGBoost	16
2.3 Long Short-Term Memory	17
2.3.1 Layers	18
2.3.2 Dropout	18
2.3.3 Preprocessing	19
Chapter 3: Results	21
3.1 Trading Strategies	22
3.2 Performance	24
3.3 Performance Net Fees	28
Discussion	29
3.4 Limitations	29
3.5 Conclusion	29
3.6 Future Work	29
Appendix A: Appendix A	31
A.1 Model Features	31
A.2 QQ Plot of Log Returns	33
A.3 Returns Net Fees	34
References	35

List of Tables

List of Figures

1.1	Hourly Tweet Counts	8
1.2	Hourly Gtrends ratios	11
1.3	12
1.4	Correlations with Log Returns	13
2.1	Distribution of Log Returns	16
2.2	17
3.1	Zcash Price Across Training, Validation and Testing Sets	21
3.2	XGBoost Predictions	23
3.3	LSTM Predictions	24
3.4	Per Trade PNL	25
3.5	XGBoost Strategies PNL	26
3.6	LSTM Strategies PNL	27
A.1	33
A.2	34

Abstract

This is a placeholder

Dedication

You can have a dedication here if you wish.

This is a placeholder for the abstract.

Introduction

Since Bitcoin's introduction in 2008, Blockchain has been touted as a revolutionary technology capable of rapidly modernizing the entire financial system in part by enabling the spread of cryptocurrencies. While these lofty goals have yet to be realized, the rapid ascension of Bitcoin and many other cryptocurrencies have made many extremely wealthy. Hobbyist investors hoping to replicate these same fortunes, have flocked to cryptocurrencies in search of high volatility investment vehicles with significant upside.

It is not a stretch to claim that these marketplaces have displayed many characteristics more synonymous with gambling than other regulated financial products. Market manipulation, specifically pump and dump schemes are rampant and relatively unchecked with Reddit and other forums dictating large market movements(Tao Li, Shin, & Wang, 2019). A recent study also found strong correlations between cryptocurrency trading and both high risk stock trading and gambling, the latter of which has been shown to succumb to human biases(Mills & Nower, 2019, Vernon (2003)).

A core feature of cryptocurrency markets is the lack of easily calculable intrinsic value in the underlying products. The value of a cryptocurrency is largely determined by its popularity and usage, leading to a feedback loop as value and popularity build upon each other. Without clear bounds of reasonable valuations, these currencies make large swings often with no underlying changes in the core products or macroeconomic landscape. For the statistically inclined, this has provided an interesting opportunity to attempt to model these seemingly random fluctuations in both price and sentiment. JP Koning describes the problem as attempting to "anticipate what average opinion expects average opinion to be", no simple task(2019).

While sentiment indicators have been used to predict equities for decades, they seem particularly applicable in the prediction of cryptocurrencies due to their derivation of value. Researchers have shown that forum comments, Wikipedia and Google search volumes, and Twitter comments can be utilized to predict a variety of cryptocurrencies with varying levels of accuracy(Kristoufek, 2013,Tianyu Li, Chamrajnagar, Fong, Rizik, & Fu (2019), Y. Kim et al. (2001)). Long short term memory (LSTM) neural nets, extreme gradient boosted trees (XGBoost), and vector autoregression have been shown to have success in this context.

However, much of the academic research in this space has focused on Bitcoin and the use of a single class of predictors. Many researchers have either shown correlations between predictions and prices or success in binary prediction, but have failed to propose profitable strategies that consistently outperform the market(Tianyu Li et al.,

2019, Colianni, Rosales, & Signorotti (2015)).

In this paper, I predict price movements in altcoin, Zcash. Zcash originated as a fork from the Bitcoin blockchain in 2016 albeit with a different name: Zerocoin protocol. Zcash provides the option for users to shield their identities and amount during a transaction, through zero-knowledge cryptography. Although it has gained recognition through its privacy features, as of May 2018 less than one percent of transactions were fully shielded, supporting the thesis that those trading Zcash are apathetic towards the specifics of the technology(Floyd, 2018). I chose Zcash, ticker ZEC, because it has moderately high volume, twenty-second in market capitalization among cryptocurrencies, has stable but developing technology, and has a unique and distinguishable name/ticker.

I use search volume, buy/sell volume, and twitter text sentiment combined with XGBoost and LSTM Neural nets to predict hourly price movements in Zcash. I propose strategies that are shown, through the use of walk-forward testing, to return over 200 percent net fees in a span of approximately one hundred days. My approach adds to the current state of academic research by utilizing this combination of features to propose strategies that return substantial profits in an altcoin market. Furthermore, these strategies and models operate at the hourly level, whereas previous research has focused on daily returns.

In Section 1, I use advanced methods to estimate tweet sentiment based on hourly data with tweets from suspected bots removed. I normalize Google search volume for Zcash by benchmarking it to the search volume of the largest cryptocurrency, Bitcoin. In section 2, I detail the XGBoost and LSTM models I use as well as my decision to regress on log returns. In section 3 I use walk-forward testing to estimate the theoretical profits of trading strategies based on these predictions. I dig into the models in an attempt to parse out how sentiment and volume may interact to drive price movements in certain situations.

In conclusion, I discuss what these results may demonstrate regarding market dynamics. I also examine the feasibility of implementation given the current exchange landscape, including liquidity and fees.

Literature Review

Early work demonstrated the predictive power of twitter sentiment on large-cap stocks(Tushar & Srivastava, 2012). Oliveira et al. used a Kalman filter to combine twitter sentiment with longer-term popular survey sentiment indicators(2017).

Recently this research has been extended into cryptocurrency, arguably a much better application, with academics demonstrating the predictive power of Twitter sentiment and Wikipedia and Google search volume on cryptocurrency prices. Garcia and Schweitzer provided proof of concept using Twitter sentiment and volume to predict Bitcoin prices(2015). They made additional use of technical market features and applied a multi-dimensional model of vector autoregression to achieve a Sharpe ratio of 1.77. They found that higher polarization was correlated with rises in Bitcoin.

Tianyu Li et al. demonstrated the effectiveness of Twitter sentiment in predicting the altcoin, Zclassic(Tianyu Li et al., 2019). They determined that 3 hours was the timeframe over which sentiment is fully absorbed into the price, and after extensive cross-validation, they had the most success using extreme gradient boosting. They made use of an additional feature that more heavily weighted retweets due to their network effect.

Kim et al. displayed a similar effect with user comments on relevant cryptocurrency online communities, displaying predictive power on both price and volume for several of the largest cryptocurrencies(2001). They utilized sets of keywords to determine relatedness to certain concepts in addition to the direction of sentiment.

ElBahrawy et al. displayed significant theoretical returns on a Wikipedia based trading strategy(2019). Kristoufek paired the Wikipedia data with Google trends to develop a multi-feature trading strategy using vector autoregression and vector error-correction(2013). He demonstrated that when prices are high increased interest tends to push prices higher, and vice-versa. This behavior may explain both the high volatility and frequencies of bubbles in the crypto space. This work was continued by Alexander Dickerson who used a set of Wikipedia and Google search frequencies to build trading strategies(2018). While they appear to be extremely profitable due to testing throughout the bitcoin boom they are not able to adequately predict the subsequent crash. As the cryptocurrency space has experienced many long term swings, it is important that these algorithms are robust enough to remain accurate in the event of novel changes in the dynamics of the market. This is an issue across much of the research in this space as studies often have too small a scope to measure realistic applicability or they predict across the entirety of the cryptocurrency era in which prices, for the most part, have consistently risen. The latter allows models to

overfit and display high theoretical profitability.

Outside of the sentiment/social information space, Wei demonstrated success using solely technical market information and an LSTM neural net(2018). Using a sliding window transformation, he exhibited profits while trading on a daily timescale producing an F1 score of almost 60% when making binary predictions.

It is important to note that while all of this proof of concept research has demonstrated significant profits, few have taken into account transaction fees or the price of liquidity. It is unclear whether these strategies would be profitable if implemented in the real world, although many show high enough returns to infer profitability. The combination of transaction fees and spreads can be in the range of 20 to 100 basis points, therefore profitable strategies require a significant edge per trade. In practice, various tactics can be used to reduce transaction costs of different strategies, but the impact of these tactics will vary by strategy.

Chapter 1

Data

Data was collected in weekly increments from October 20th to March 17th, accumulating in over 3,500 data points.

1.1 Twitter

Twitter is one of the largest social media platforms, distinct in its focus on brief text. It has a large cryptocurrency community with a combination of individual traders and key moguls who dictate the trading flow of thousands of followers.

The Twitter API allows one to pull all tweets containing certain keywords. I collected all tweets in English containing the terms ‘zcash’ or ‘zec’. With over 130 thousand relevant tweets during the specified time period, there are on average 37.5 tweets every hour before accounting for bots or other confounding factors.

To remain consistent with the other features, I aggregated the Twitter data by the hour. As shown in Figure 1.1, while there is a daily trend in the number of tweets every hour the autocorrelation is not extreme enough to require adjustment. I chose to explore the mean sentiment, a favorite weighted sentiment, and a retweet weighted sentiment as well as the tweet count. Testing showed that the mean sentiment, retweet weighted sentiment, and tweet count, were instrumental to the model, while the favorite weighted sentiment seemed to add noise and caused the model to overfit.




Figure 1.1: Hourly Tweet Counts

1.1.1 Tweet Removal

While I previously described the unique name as being one of the deciding factors in choosing Zcash, a hurdle arose as the string ‘zec’ appears both in Zechariah, a Hebrew prophet, and is an acronym for the Zimbabwe Electoral Commission. A nonnegligible portion of the tweets collected, approximately 2-3%, referenced one of these two topics. Therefore, I removed all tweets containing ‘Zechariah’ as well as common terms appearing in relation to the election: ‘election’, ‘Zimbabwe’, ‘Botswana’, etc..

Twitter has taken aggressive measures to remove bots from its platform. However, a recent study estimated that 9 to 15% of accounts are bots while another study found that two-thirds of shared links are posted by bots (Wojcik, n.d., Varol, Ferrara, Davis, Menczer, & Flammini (2017)). This may be even higher in the cryptocurrency space due to the potential profits if one can swing public opinion. I use Michael Kearney’s ‘tweetbotornot’ package to eliminate this noise by classifying and removing bots. Michael Kearney’s package applies an extreme gradient boosting model to a user’s tweets and bio to predict whether or not they are a bot with over 93% accuracy. Due to the high number of bots and the large number of tweets that they post, I do not expect these accuracy rates to transfer perfectly onto my dataset. I manually tested 30 users that were classified as bots and 30 users classified as non-bots, the results are shown in Table 1.1.1. An unknown refers to an account that was removed, which we

can speculate would often be bots. This test used a .7 threshold that was selected in the validation phase. My own distinctions are not objective, so blind testing was performed to eliminate partial bias from these results. The results show statistically significant differences between the groups. Out of the 15 bots, most were classified correctly, however there was still a significant portion of false positives. While this is concerning, it seems likely that for modeling purposes it is more important to remove the systematic bias of bots than to capture the full extent of the authentic signal. This was supported empirically as it was advantageous to remove a larger portion of bots at the expense of removing some authentic accounts.

Bot Model Accuracy

True

False

Unknown

Bot

14

4

12

Non-Bot

26

1

3

Overall slightly over half of the users and half of the tweets were removed. After removing tweets not related to the cryptocurrency Zcash and those posted by potential bots, the dataset contains approximately 57 thousand tweets for an average of 16 tweets per hour.

1.1.2 Sentiment Analysis

R has multiple packages that perform sentiment analysis on texts. I implemented two of the most popular libraries The ‘SentimentAnalysis’ package allows the use of many different dictionaries, however, I chose the two most suitable to my goals: Henry’s finance specific dictionary and the QDAP dictionary put together by Tyler Rinker. ‘SentimentR’ similarly utilizes the QDAP dictionary, however it also takes into account valence shifters such as negators to improve performance over a simple dictionary lookup while still maintaining speed. To assess performance, I randomly sampled 100 tweets and cross-checked their predictions using my own classifications as the ground truth. The results are shown in Table 1.1.2. Tweets were classified into three groups: positive, neutral, and negative. ‘Sentimentr’ had the best performance when looking at weighted F1 scores, which given the class imbalance is a better measurement of accuracy.

$$Weighted\ F1 = \sum_{i=1}^{numclasses} 2 * \frac{(precision_{class_i} + recall_{class_i})}{(precision_{class_i} * recall_{class_i})} \quad (1.1)$$

(1.2)

Sentiment Package Accuracy

Henry's Financial

QDAP

SentimentR

F1

0.55

0.62

0.62

Weighted F1

0.62

0.63

0.70

The confusion matrix using 'SentimentR' is in Table 1.1.2. It is important to note that most of the errors occur through neutral samples being labeled as positive or negative and vice-versa. There are relatively few occurrences of negative/positive samples being classified as the opposite. One example of a negative outlook tweet that was labeled as positive is the following: "All privacy coins are held together with duct tape + paperclips, in particular, all XMR and ZEC forks. Epic work by...". This is a tweet that even the most sophisticated NLP algorithm would struggle with due to the ambiguity of the metaphor and the positive connotation of "epic".

SentimentR Confusion Matrix

Obs. Negative

Obs. Neutral

Obs. Positive

Negative

6

8

9

Neutral

0

26

8

Positive

3

14

26

1.2 Google Search Volume

Google trends display normalized search volume for a given keyword. I collected trends data for the terms 'zec', 'zcash', and 'bitcoin'. I use the ratio of 'zec' and 'zcash' search volumes to 'bitcoin' in order to normalize for general trends in the interest in cryptocurrency and also included the raw 'zcash' value. Figure 1.2 illustrates that the hourly ratios do not display autocorrelation to an extent that requires correction.




Figure 1.2: Hourly Gtrends ratios

1.3 Price and Volume

Cryptocompare uses a sophisticated aggregation method to accurately assess real prices for cryptocurrencies given the multitude of markets with different liquidity and regulations. They use a volume-weighted average with an additional decay based on the time since the most recent trade and outlier detection.

Volume to and volume from correspond to the volume in and out of a certain currency in the pair. I use the conversion rate of ZEC to USD. Interestingly as shown in ??, as price increased in January and early February the volume initially increased dramatically, but then tailed off.

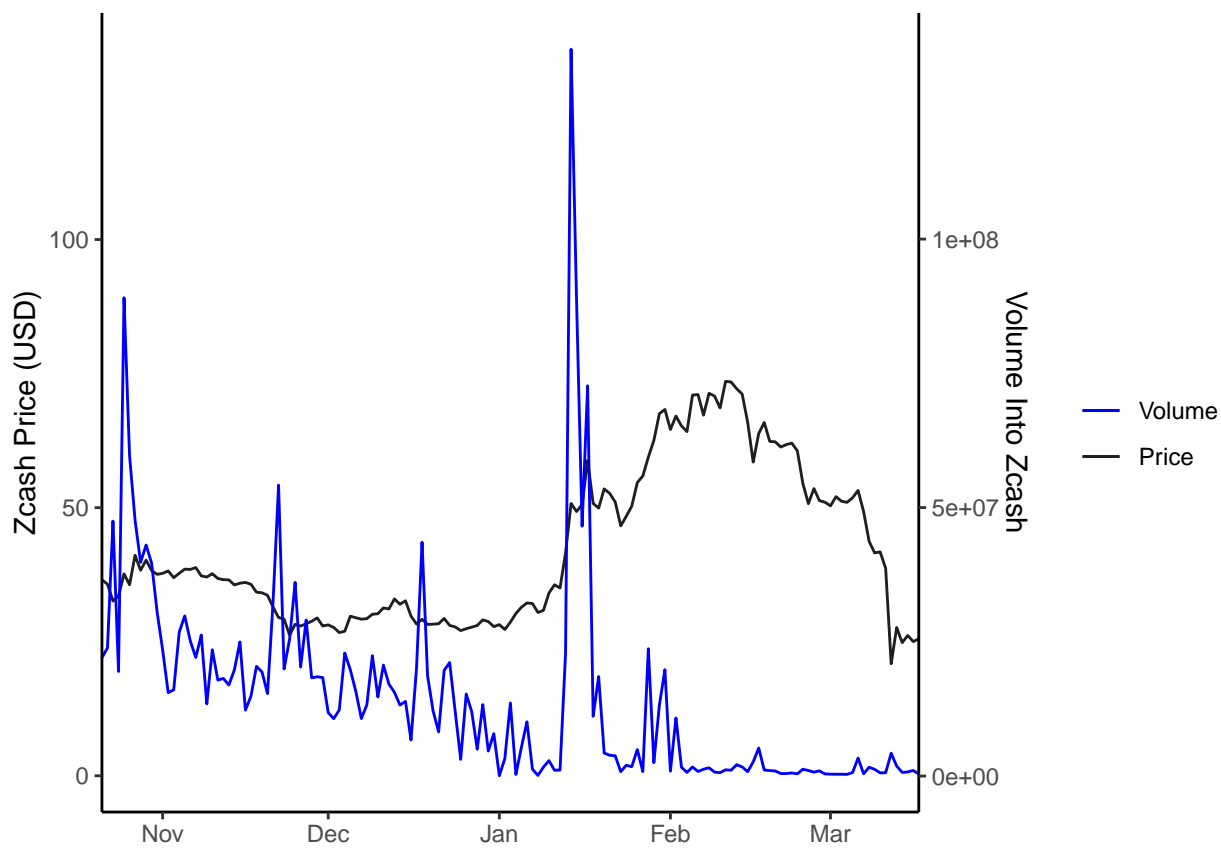


Figure 1.3

1.4 Exploratory Data Analysis

Correlations between the model features and four hour log returns are shown in Figure 1.4.

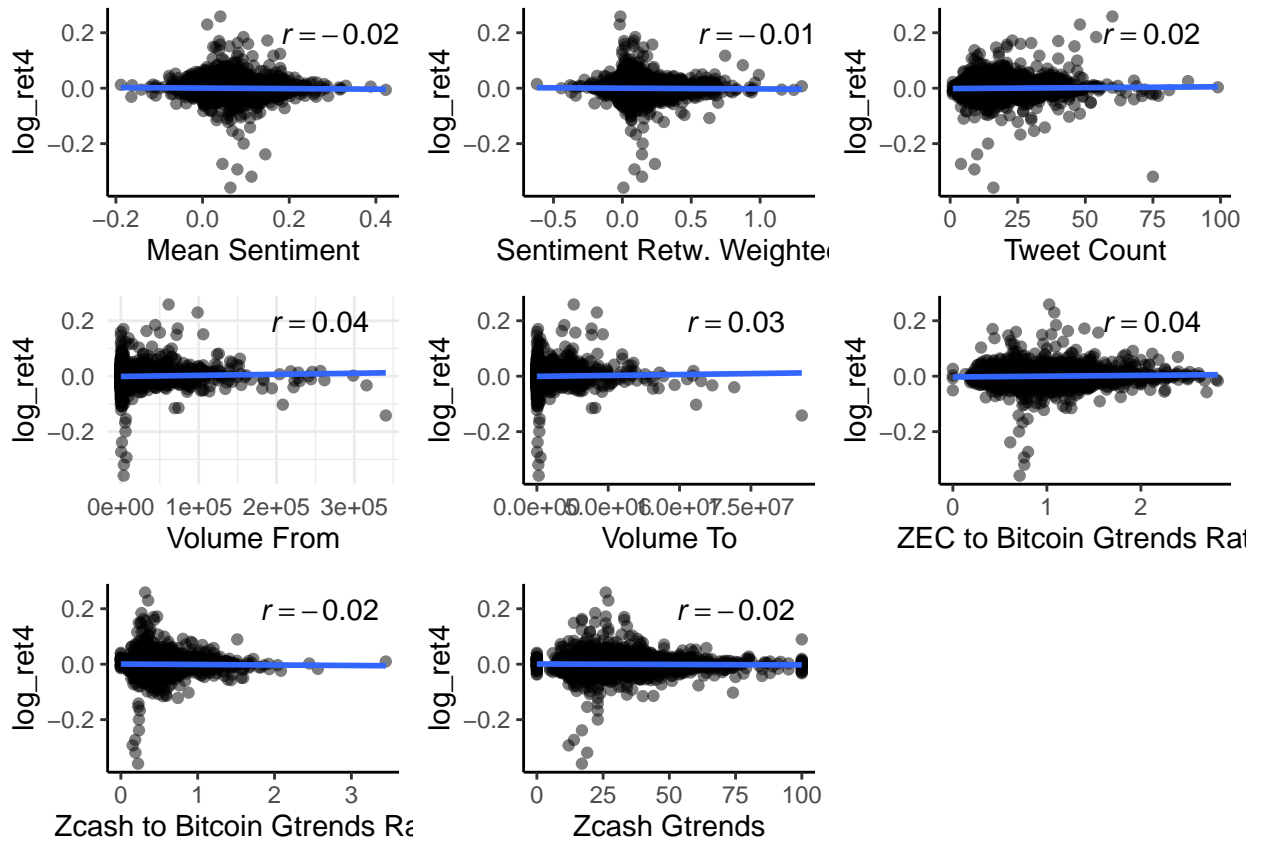


Figure 1.4: Correlations with Log Returns

Chapter 2

Methodology

2.1 Log Returns

I regress on log returns with a specified lag of 4 hours that was determined in the validation phase. This is not necessarily equivalent to the time it takes for the full effect of the sentiment to be realized, but may be an optimization of the tradeoff between signal and noise.

$$\log return_i = \log\left(\frac{p(i)}{p(i-lag)}\right) \quad (2.1)$$

$$(2.2)$$

This provides multiple advantages over price or raw returns (“Why log returns,” 2012).

1. Under the assumption that prices follow a log normal distribution than log returns are normally distributed
2. When returns are small log-returns are close in value
3. Compound returns are normally distributed as the sum of normally distributed log returns

The distribution of log returns for Zcash is shown in Figure 2.1. It should be noted that log returns do not fully adhere to a normal distribution due to the frequent multiple standard deviation occurrences¹.

¹For QQPlot see appendix




Figure 2.1: Distribution of Log Returns

2.2 XGBoost

Extreme Gradient Boosted Trees, XGBoost, is a decision tree based ensemble learning method that employs boosting to efficiently yield high accuracy predictions. Tree ensemble methods aggregate the predictions of weak decision trees constantly updating residuals to improve predictions. Assuming K regression trees, our prediction for point y_i is the sum of the feature set \mathbf{x}_i regressed on each individual tree f_k .

$$\hat{y}_i = \phi(\mathbf{x}_i) = \sum_{k=1}^K f_k(\mathbf{x}_i), \quad f_k \in F \quad (2.3)$$

$$(2.4)$$

We minimize the following objective function where Ω is a regularization term that penalizes more complex models. T is the number of leaves and w is the coefficient at each node. In my model l is taken to be the mean squared error.

$$L(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k) \quad (2.5)$$

$$\Omega(f) = \gamma T + 1/2\lambda \|w\| \quad (2.6)$$

$$(2.7)$$

The model is trained in an iterative manner. If we let $\hat{y}_i^{(t)}$ be the prediction for point y_i at the t iteration than we design f_t to minimize the following function.

$$L(\phi) = \sum_i^n [l(\hat{y}_i^{(t-1)} + f_t(\mathbf{x}_i), y_i)] + \Omega(f_t) \quad (2.8)$$

$$(2.9)$$

To achieve this I employ the following algorithm designed by Chen and Guestrin (2016).

Algorithm 1: Exact Greedy Algorithm for Split Finding

Input: I , instance set of current node

Input: d , feature dimension

$gain \leftarrow 0$

$G \leftarrow \sum_{i \in I} g_i, H \leftarrow \sum_{i \in I} h_i$

for $k = 1$ **to** m **do**

$G_L \leftarrow 0, H_L \leftarrow 0$

for j **in** $sorted(I, \text{by } \mathbf{x}_{j_k})$ **do**

$G_L \leftarrow G_L + g_j, H_L \leftarrow H_L + h_j$

$G_R \leftarrow G - G_L, H_R \leftarrow H - H_L$

$score \leftarrow \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$

end

end

Output: Split with max score

Figure 2.2

2.3 Long Short-Term Memory

Long Short-Term Memory is a type of recurrent neural network that due to the specific architecture is able to better discern longer term patterns, avoiding the vanishing gradient problem (Hochreiter & Schmidhuber, 1997). It has been used in many contexts, making notable gains in speech recognition, language translation, and general time series prediction (Beaufays, 2015).

The basic architecture of a memory cell in an LSTM network consists of an input gate, forget gate, cell state, and output gate. The input gate I_t , forget gate F_t , and output gate O_t for memory cell t are linear combinations of the previous hidden state H_{t-1} and current input X_t passed through the sigmoid function. I denote the weight

vectors as U and W and \odot and $+$ denote elementwise multiplication and addition.

$$I_t = \sigma(X_t \odot W_i + H_{t-1} \odot U_i) \quad (2.10)$$

$$F_t = \sigma(X_t \odot W_f + H_{t-1} \odot U_f) \quad (2.11)$$

$$O_t = \sigma(X_t \odot W_o + H_{t-1} \odot U_o) \quad (2.12)$$

$$(2.13)$$

The cell state C_t is then a combination of the previous cell state filtered by the forget gate and the input gate, which determines the values to be updated, elementwise multiplied by another linear combination of the hidden state and current input scaled by the tanh function(Sanjeevi, 2018).

$$c_t = \tanh(X_t \odot W_c + H_{t-1} \odot U_c) \quad (2.14)$$

$$C_t = \sigma(C_{t-1} \odot F_t + I_t \odot c_t) \quad (2.15)$$

$$(2.16)$$

The hidden state is the current state scaled by the tanh function weighted by the output gate.

$$H_t = \tanh(C_t) \odot O_t \quad (2.17)$$

$$(2.18)$$

2.3.1 Layers

One challenge is determining the optimal number of hidden neurons and hidden layers to avoid overfitting while still capturing potentially complex interactions. From Jeff Heaton, one hidden layer can “approximate any function that contains a continuous mapping from one finite space to another”(2008). While I also tested using bidirectional layers and multiple hidden layers, I achieved the greatest accuracy through a single hidden layer.

2.3.2 Dropout

Dropout is a common technique to reduce overfitting and make individual hidden neurons more robust. Dropout refers to the random dropping of individual neurons in the training phase. In an LSTM network, this can take many different forms. After testing, I determined a 20 percent recurrent dropout in the LSTM layer combined with a 40 percent dropout layer was the most effective. Recurrent dropout in the context of LSTM refers to dropping d during the update of the current state(G, 2018).

$$C_t = \sigma(C_{t-1} \odot F_t + I_t \odot d(c_t)) \quad (2.19)$$

$$(2.20)$$

I then added the additional, stronger dropout layer on top of the LSTM layer. I thus avoided a common issue in LSTM dropout in which important past information is dropped, reducing the ability of the model to identify long term patterns.

2.3.3 Preprocessing

Unlike tree based models which are insensitive to the range of data, LSTM perform best when the data is normalized. I normalized each feature with the minimum and maximum of the training data².

In addition, I used a sliding window lookback as seen in (Wei, 2018). I found through validation that a lookback of 4 hours combined with 8 hidden units had the most success.

²This avoids potential bias if the test data impacts the scaling

Chapter 3

Results

The hyperparameters for the models were tuned based on cross validation performed on the first 30 percent of the timeseries (1,060 data points) as shown in Figure 3.1. It should be noted that the test period was a particularly volatile interval for Zcash, which doubled before crashing back to below initial levels.




Figure 3.1: Zcash Price Across Training, Validation and Testing Sets

Performance was measured using anchored¹ walk-forward testing, as seen in (Żbikowski, 2014), in which one iteratively predicts test data point i using a model

¹Unanchored walk-forward testing uses a sliding window of training data to remove outdated datapoints. This is unnecessary given the scope of training data.

trained on data points $(1 : i - 1)^2$. This avoids the issue of look-ahead bias that occurs when standard cross validation techniques are used on time series data, and allows us to best simulate a live trading algorithm. Walk-forward testing provides a more realistic measure of model performance as the model constantly updates to include new market conditions and tests the stability of the model to new data.

It should be noted that while the XGBoost model was retrained on the insertion of each new data point, due to time constraints and to reduce overfitting, the LSTM model was trained for one epoch every 50 new data points, approximately two days of data. Both models were trained to minimize root mean squared error.

3.1 Trading Strategies

The following strategies described in Table 3.1 were tested for both XGBoost and LSTM predictions, using the returns of holding Zcash as a baseline strategy. One advantage of using prediction percentiles instead of raw values to inform strategies is that it improves the balance between the number of buy/sell trades, reducing the impact of directional bias. I tested both an exponential decay and weighting by prediction value, however these approaches did not yield significantly better results. The percentiles of predictions were calculated for prediction i on the set of predictions $(1 : i)$ with the notable exception being that the first ten percent of test data points were grouped together.

Strategy Descriptions

Strategies

Descriptions

Hold

Return of holding Zcash

All

Trade on every increment, direction in comparison to mean prediction

Top 20

Trade on the top and bottom 20 percent of predictions

Top 10

Trade on the top and bottom 10 percent of predictions

The predictions for the XGBoost and LSTM models are displayed below in Figures 3.2 and 3.3. The points are subsetting based on the Top 20 trading strategy. The XGBoost predictions are more evenly distributed throughout the time series, whereas the LSTM predictions follow clear trends.

²As the model fits on four hour lagged returns, points $(i - 4 : i - 1)$ will be correlated with point i ; therefore, this process was modified to train on points $(1 : i - 5)$.




Figure 3.2: XGBoost Predictions

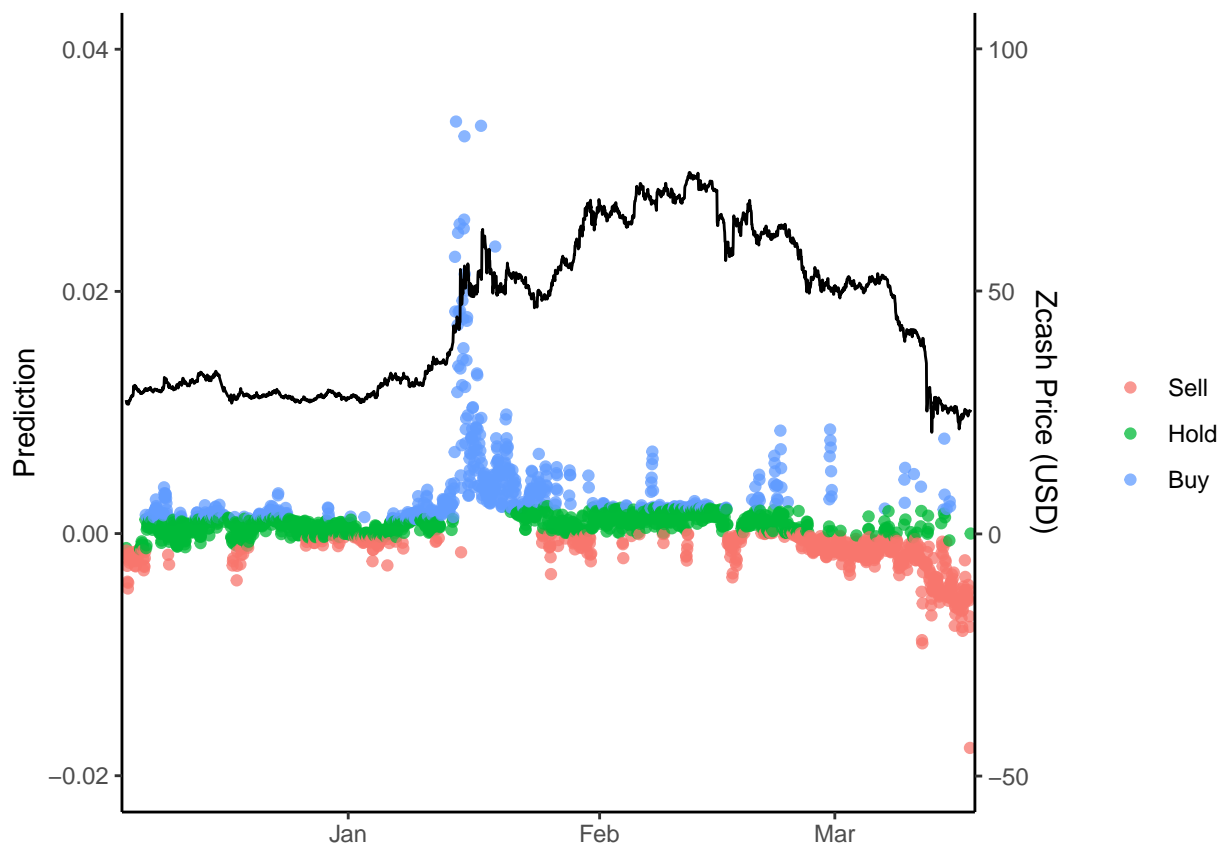


Figure 3.3: LSTM Predictions

3.2 Performance

While the models achieved the best performance when trained on 4 hour log returns, we actually see that the performance accumulates past 4 hours. We see a monotonic increase in mean performance across all the strategies in Figure 3.4 up to 5 hours after the trade. The LSTM and XGBoost Top 20 strategies return over a third of a percent per trade.

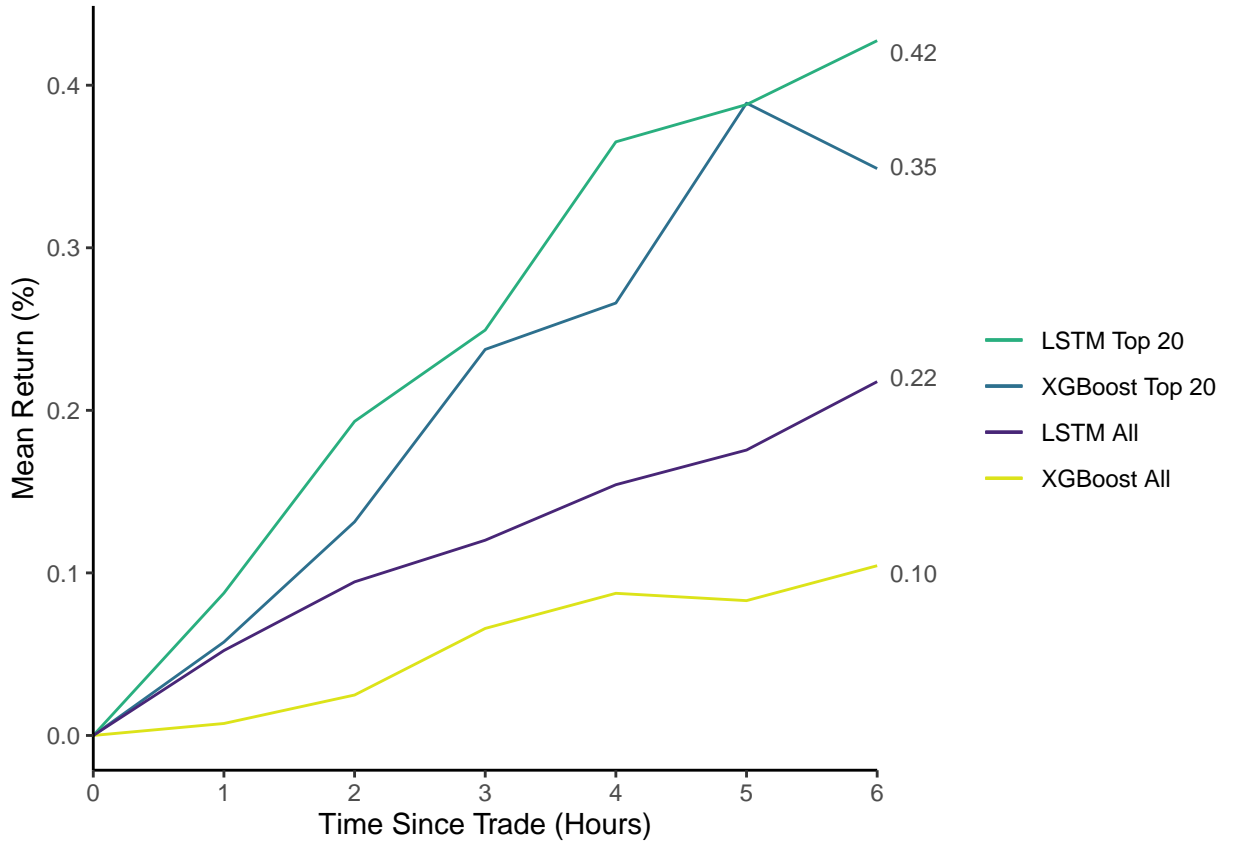


Figure 3.4: Per Trade PNL

Per trade returns are slightly misleading as many signals overlap, resulting in actual profits that are much lower than pure sum of mean returns. Figures 3.5 and 3.6 display the theoretical returns of the strategies over the testing period. In the implementation of these strategies it is assumed that new signals override old predictions and after 5 hours of no prediction the position is closed.




Figure 3.5: XGBoost Strategies PNL

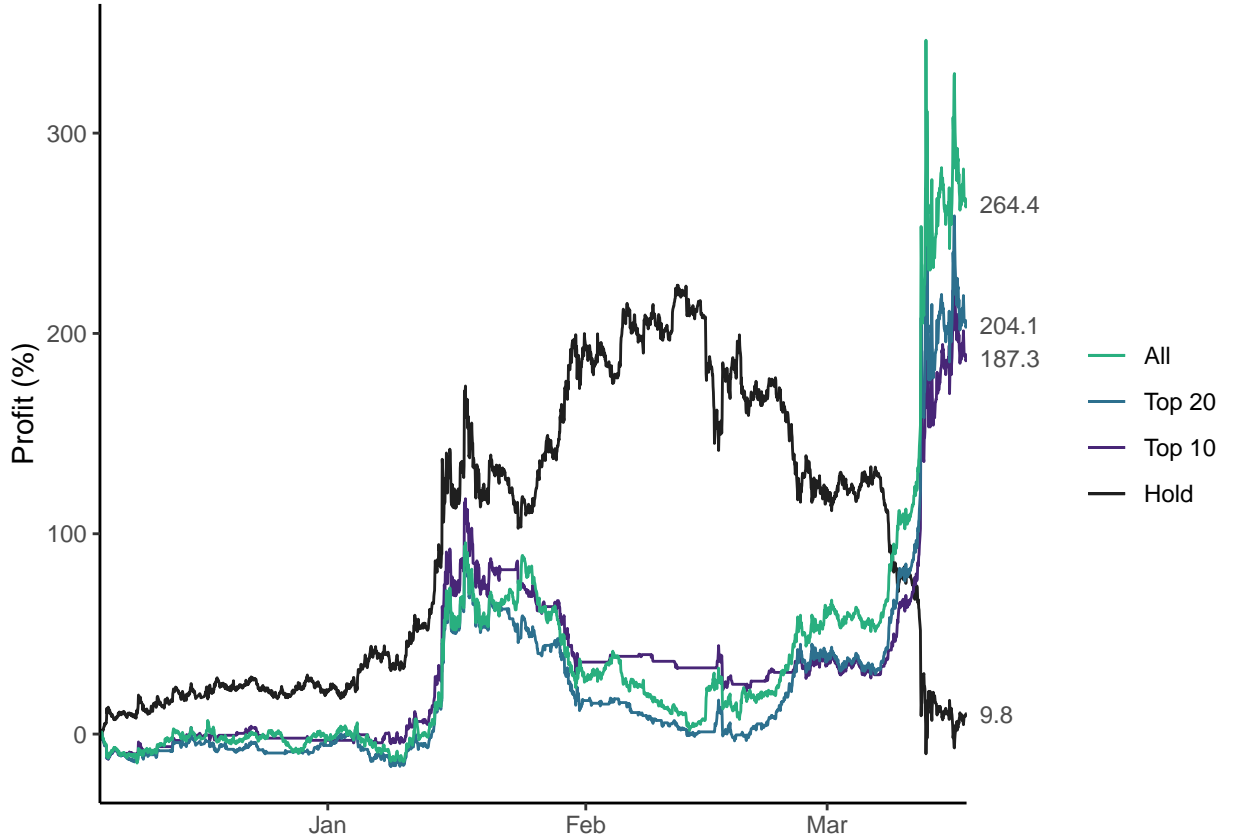


Figure 3.6: LSTM Strategies PNL

All three LSTM strategies outperform the most profitable XGBoost strategy. The LSTM All and Top 20 strategies both return over 200 percent in just over 100 days while the XGBoost Top 20 strategy also produces sizeable returns of 94 percent. Every strategy except the XGBoost Top 10 outperforms the market during this time period.

While by no means a perfect metric, the Sharpe ratio is a strong proxy for risk adjusted returns. It is the excess expected return compared to the risk free rate R_f divided by the standard deviation of returns.

$$Sharpe = \frac{E(Return - R_f)}{sd(Return - R_f)} \quad (3.1)$$

$$(3.2)$$

Using a risk free rate of 2%, Table 3.2 displays the Sharpe ratio for the following strategies. This is calculated with daily returns, and then scaled by the square root of days in a calendar year.

Sharpe Ratios
 XGBoost Top 20
 LSTM Top 20
 LSTM All
 Sharpe Ratio
 3.11

2.67

2.67

While these strategies display strong Sharpe ratios, they are depressed by the large volatility in the test data. Although the XGBoost 20 strategy was less profitable it has a higher Sharpe ratio because of less volatility in its returns.

3.3 Performance Net Fees

Transaction costs represent a considerable expense in the cryptocurrency market. As a point of reference, Binance, one of the largest cryptocurrency exchanges, charges transaction fees of between 1.5 to 7.5 basis points³. While spreads are generally fairly negligible on liquid markets, during times of high volatility they can represent an additional 5-10 basis points. Table 3.3 displays the strategy profits assuming different fee structures. It is clear that even with substantial transaction fees these strategies remain profitable. However, when taking transaction fees into account it would be advantageous to only trade on the most profitable subset of signals.

Percent Returns Net Fees

15 bps

5 bps

2.5 bps

XGBoost Top 20

11.5%

61.3%

76.9%

LSTM Top 20

156.7%

187.4%

195.7%

LSTM All

134.7%

214.7%

238.6%

³One hundredth of one percent

Discussion

3.4 Limitations

Though this model displays a notable edge in testing, it is important to recognize potential limitations that might make it unprofitable in practice.

To begin, there may be a mismatch between the prices in my dataset and the price point of actionable liquidity. As some of these exchanges operate in only a subset of countries, it may be the case that I have a theoretical edge, but in practice the trades would not be possible at the price points displayed. This might be the case if aggregated price points oscillate in weightings between different exchanges.

Secondly, my edge may be real but simply not account for black swan events. While I may have an on average profitable strategy, the risk adjusted rate of returns are poor.

Lastly, there is the possibility that these returns are due to pure chance. While this is a large dataset and we see strong average returns, a few small timeframes account for most of the price movement. These profits could be attributable to a few correct predictions or a few associations that only by accident work in this context. Back testing is a tedious process in which many false associations may be found.

3.5 Conclusion

The goal of this thesis is to demonstrate strategies that produce substantial alpha in an altcoin market. While I am cognizant of potential limitations, it seems likely that these strategies capture legitimate alpha by parsing out signal from a combination of features. Other studies have shown similar predictive power in Bitcoin and here I add proof of concept that models using both sentiment and volume based indicators can be similarly accurate in a smaller market with a fraction of the data.

3.6 Future Work

Additional research can improve upon the sentiment and bot classification algorithms used in this study. These algorithms can be tuned specifically to the context of cryptocurrency to achieve more accurate categorization. Furthermore, additional research should explore hedging strategies that would help reduce the variance of returns, therefore improving the risk adjusted returns.

Appendix A

Appendix A

A.1 Model Features

Zcash Gtrend

Zcash Gtrend Ratio

ZEC Gtrend Ratio

Volume From

Volume To

Tweet Count

Mean Sentiment

Sentiment Retw. Weighted

33

1.100000

2.133333

461959.8

12541.98

12

0.0573850

0.0000053

46

1.586207

1.551724

516757.6

14040.55

7

0.1481721

0.0833144

40

1.481481

2.296296

910020.4

24811.14

9
0.0196764
0.0000000
35
1.206897
2.620690
943465.4
25714.64
5
0.0097706
0.0908355
26
0.962963
2.518518
879135.5
24370.13
14
0.0735963
0.1326699

A.2 QQ Plot of Log Returns




Figure A.1

A.3 Returns Net Fees




Figure A.2

References

- Beaufays, F. (2015, August). Google ai blog. *Google AI Blog*. Retrieved from ai.googleblog.com/2015/08/the-neural-networks-behind-google-voice.html
- Chen, T., & Guestrin, C. (2016). Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining* (pp. 785–794).
- Colianni, S., Rosales, S., & Signorotti, M. (2015). Algorithmic trading of cryptocurrency based on twitter sentiment analysis.
- Dickerson, A. (2018). Algorithmic trading of bitcoin using wikipedia and google search volume. Retrieved from <http://dx.doi.org/10.2139/ssrn.3177738>
- ElBahrawy, A., Alessandretti, L., & Baronchelli, A. (2019). Wikipedia and cryptocurrencies: Interplay between collective attention and market performance.
- Floyd, D. (2018). Zcash privacy weakened by certain behaviors, researchers say. Retrieved from www.coindesk.com/zcash-privacy-weakened-by-certain-behaviors-researchers-say
- G, A. (2018). A review of dropout as applied to rnns. medium.
- Garcia, D., & Schweitzer, F. (2015). Social signals and algorithmic trading of bitcoin. *Society Open Science*, 2(9).
- Heaton, J. (2008). *Introduction to neural networks for java, 2nd edition* (2nd ed.). Heaton Research, Inc.
- Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735–1780. <http://doi.org/10.1162/neco.1997.9.8.1735>
- Kim, Y., Lee, J., Park, N., Choo, J., Kim, J., & al. (2001). When bitcoin encounters information in an online forum: Using text mining to analyse user opinions and predict value fluctuation. Retrieved from <https://doi.org/10.1371/journal.pone.0177630>
- Koning, J. (2019, November). Moneyiness. *Moneyiness*.
- Kristoufek, L. (2013). BitCoin meets google trends and wikipedia: Quantifying the

- relationship between phenomena of the internet era. *Scientific Reports*, 3(1).
- Li, T., Chamrajnagar, A., Fong, X., Rizik, N., & Fu, F. (2019). Sentiment-based prediction of alternative cryptocurrency price fluctuations using gradient boosting tree model. *Front. Phys.*
- Li, T., Shin, D., & Wang, B. (2019). Cryptocurrency pump-and-dump schemes. Retrieved from <http://dx.doi.org/10.2139/ssrn.3267041>
- Mills, D. J., & Nower, L. (2019). Preliminary findings on cryptocurrency trading among regular gamblers: A new risk for problem gambling? *Addictive Behaviors*, 92, 136–140.
- Oliveira, N., Corteza, P., & Areal, N. (2017). Analyzing stock market movements using twitter sentiment analysis. Retrieved from <http://repositorium.sdum.uminho.pt/bitstream/1822/54457/4/paper.pdf>
- Sanjeevi, M. (2018, January). Chapter 10.1: DeepNLP — lstm (long short term memory) networks with math.
- Tushar, R., & Srivastava, S. (2012). Analyzing stock market movements using twitter sentiment analysis. *International Conference on Advances in Social Networks Analysis and Mining*. Retrieved from <http://dx.doi.org/10.1109/ASONAM.2012.30>
- Varol, O., Ferrara, E., Davis, C. A., Menczer, F., & Flammini, A. (2017). Online human-bot interactions: Detection, estimation, and characterization. In *Eleventh international aaai conference on web and social media*.
- Vernon, A. M. (2003). Market efficiency and march madness: Empirical tests of point spread betting. *SSRN Electronic Journal*.
- Wei, J. (2018, December). Predicting cryptocurrency prices with machine learning. Retrieved from medium.com/datadriveninvestor/predicting-cryptocurrency-prices-with-machine-learning-1b5a711d3937
- Why log returns. (2012, November). Quantivity. Retrieved from quantivity.wordpress.com/2011/02/21/why-log-returns.
- Wojcik, S. (n.d.). 5 things to know about bots on twitter. Retrieved from <https://www.pewresearch.org/fact-tank/2018/04/09/5-things-to-know-about-bots-on-twitter/>
- Żbikowski, K. (2014). Using volume weighted support vector machines with walk forward testing and feature selection for the purpose of creating stock trading strategy. *Expert Systems with Applications*, 42. <http://doi.org/10.1016/j.eswa.2014.10.001>